

Kivi is a voice-first computing interface that bridges natural dictation with intelligent workspace execution. Existing voice tools force users to repeatedly re-explain project context, formatting preferences, and organizational acronyms across different desktop applications. Kivi solves this context loss by extracting durable semantic memory from daily speech. Acting as a high-precision co-pilot, Hey Kivi retrieves cross-session facts and user preferences to execute complex formatting and retrieval tasks—delivering personalized voice computing without storing raw clutter, making unverified assumptions, or compromising user agency.